

Databases for GenAI

Embeddings, Vector Databases & Production
Patterns

October 2025



Speaker



Maksym Lypivskyi
Head of Cloud Platforms &
AI Director

- 9 years at Ciklum driving large-scale cloud and software delivery initiatives
- 3 years specializing in AI
- Core interest: making AI systems reliable, production-ready, and business-impactful

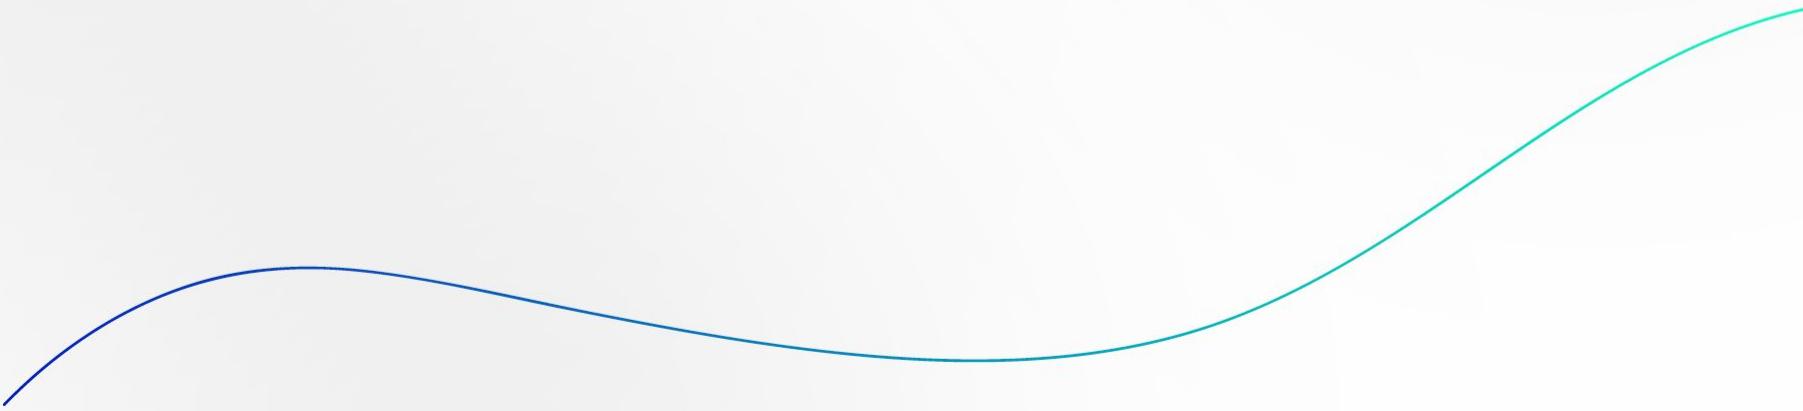
Agenda

01 Foundations

02 Advanced RAG patterns

03 Production readiness

Foundations



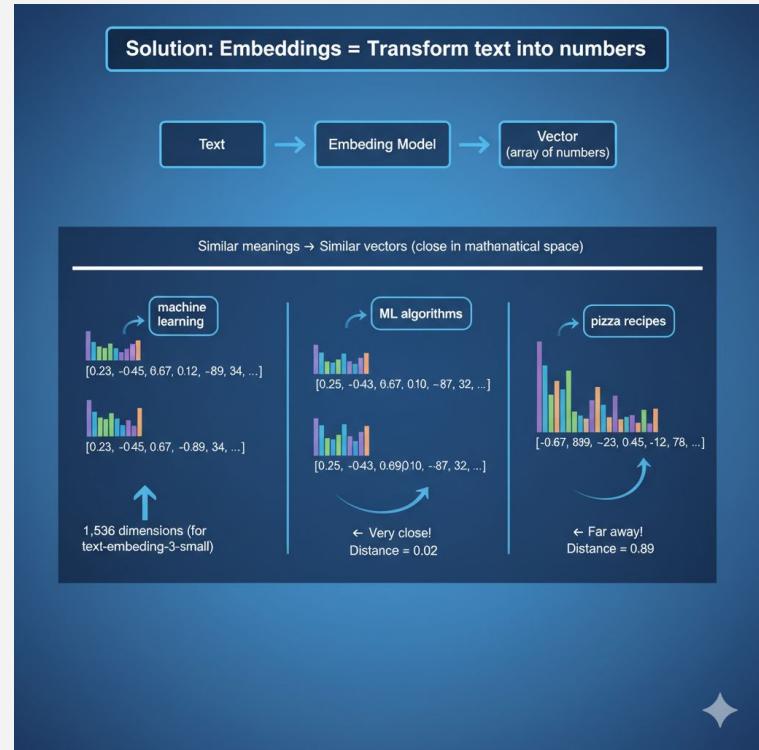
What Are Embeddings?

The Problem: Computers don't understand text meaning



"machine learning
algorithms"
"ML techniques"
"deep neural networks"

Are these similar? How
do we compute that?



How Similarity Search Works

The Vector Search Process



1. INDEXING

Documents → Embedding Model → Store vectors

Doc 1: "Python tutorial" → [0.2, -0.4, 0.6, ...] → Store

Doc 2: "Java programming" → [0.3, -0.5, 0.5, ...] → Store

Doc 3: "Cooking recipes" → [-0.8, 0.9, -0.2, ...] → Store

2. SEARCHING

User Query: "learn Python" → [0.21, -0.42, 0.58, ...]

Compute distances to ALL stored vectors:

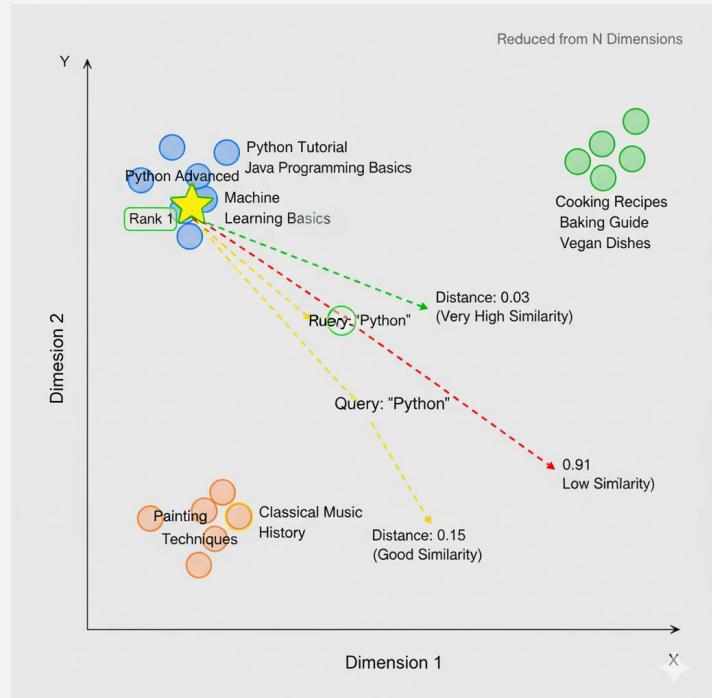
- Distance to Doc 1: 0.03 ✓ (very close!)
- Distance to Doc 2: 0.15 (somewhat close)
- Distance to Doc 3: 0.91 (far away)

Similarity Metrics

-Cosine similarity: Angle between vectors (most common)

-Euclidean distance: Straight-line distance

-Dot product: Vector multiplication



Embedding Quality Matters



EVERLAW (Legal Discovery-1.4M documents)

87%
accuracy

MINDLID

82%
3 recall

INTERACTION CO
(Email Assistant-100 emails)

21.45s
embedd

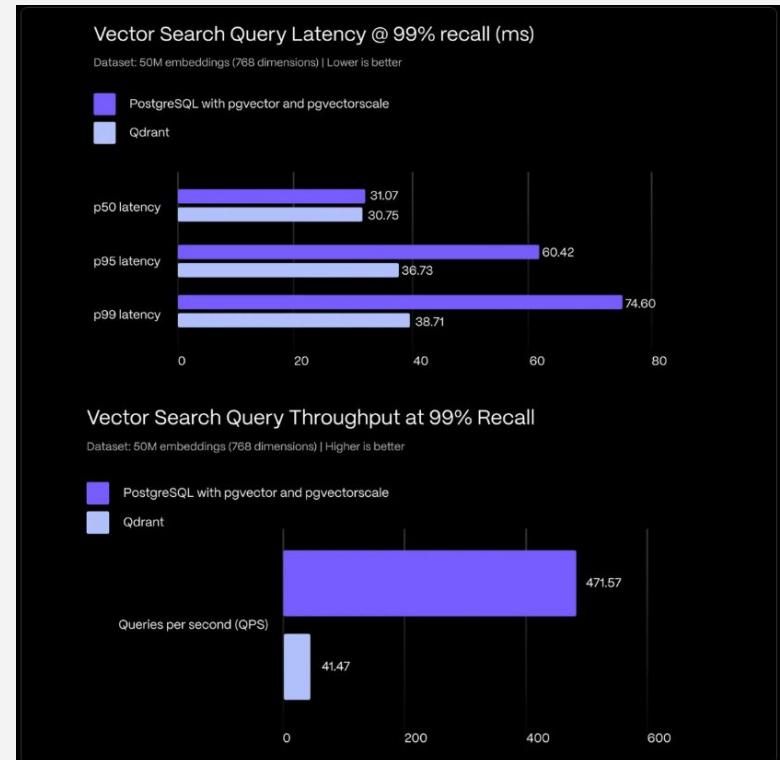
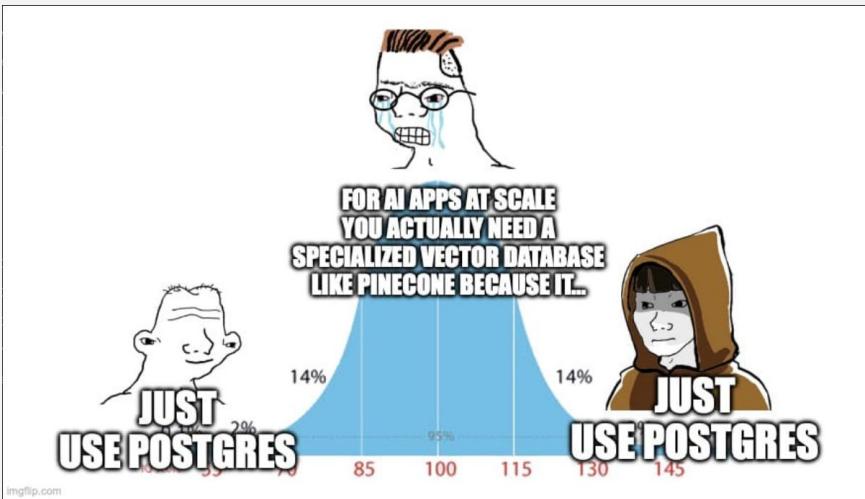
Model	Provider	MTEB	Dims	Price/1M
stella_en_1.5B_v5	NovaSearch	71.54	1024-8192	Free (OSS)
gemini-embedding-001	Google	68.32	3072	\$0.15
text-embedding-3-large	OpenAI	64.6	3072	\$0.13
text-embedding-3-small	OpenAI	62.3	1536	\$0.02

<https://huggingface.co/spaces/mteb/leaderboard>

PostgreSQL Revolution



PostgreSQL and Pgvector: Now Faster Than Pinecone, 75% Cheaper, and 100% Open Source



From MVP to Scale

Mid Size High-Performance Vector Search



Architecture & Performance

- Embedded database (runs in-process with your application)
- 112 QPS at 10M vectors
- Python and JavaScript SDKs
- Open-source (Apache 2.0 license)

Memory Optimization:

- In-memory and persistent storage modes
- Optimized for datasets under 1M vectors
- Simple distance functions (cosine, L2, inner product)

Advanced Capabilities:

- Zero-configuration setup (pip install chromadb)
- Metadata filtering
- Multiple collection support



chroma

Architecture & Performance

- Rust-based for maximum speed
- 626 QPS at 99.5% recall (1M vectors)
- Sub-3ms p95 latency
- Open-source (Apache 2.0 license)

Memory Optimization:

- Binary Quantization: 40x memory reduction
- Scalar quantization
- Product quantization

Advanced Capabilities:

- Native hybrid search (dense + sparse vectors)
- Advanced payload filtering
- Distributed architecture for horizontal scaling
- RESTful and gRPC APIs



drant

Architecture & Performance

- Built for massive scale (billions of vectors)
- 2,098 QPS at 100% recall (10M vectors)
- Sub-10ms latency on performance configurations
- Open-source (Apache 2.0 license)

Memory Optimization:

- Int8 compression: 75% memory savings
- RabitQ 1-bit quantization
- Multiple quantization strategies

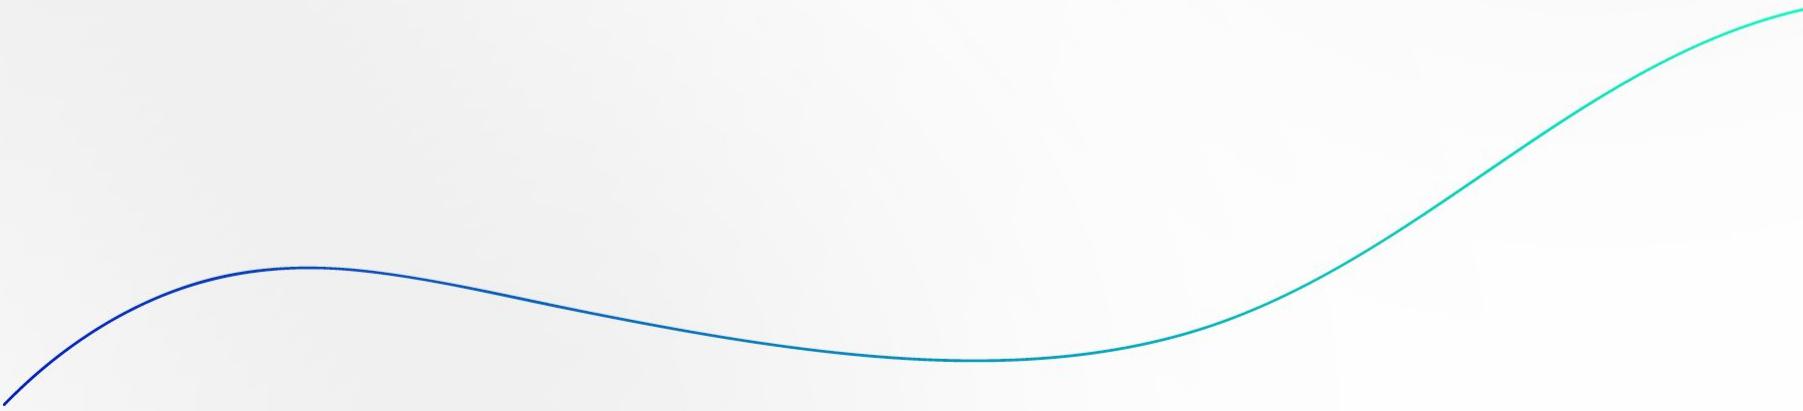
Advanced Capabilities:

- GPU acceleration support
- Up to 100,000 collections per cluster
- Distributed architecture with query node separation
- Multiple distance metrics and index types (HNSW, IVF, DiskANN)



Milvus

Advanced RAG patterns



The Multi-Hop Reasoning Problem



Complex Query:

"What are the regulatory implications of our Q3 marketing strategy for the European market?"

Traditional RAG fails



Retrieves isolated chunks:

- ✗ - "Q3 marketing strategy increased social media spend"
- ✗ - "European market regulations overview"
- ✗ - "GDPR compliance requirements"

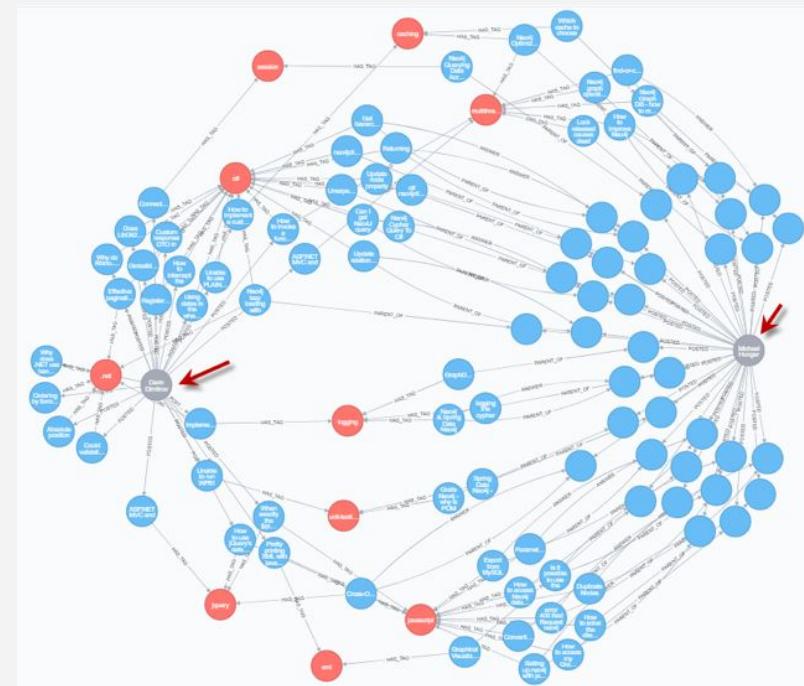
Problem: Can't connect the dots!

Marketing → Budget → Compliance → GDPR → Europe

This is where we need Graph RAG



GraphRAG Solution



The PDF Processing Problem



Traditional RAG Pipeline for PDFs

- ✗ OCR → Errors with complex layouts
- ✗ Layout detection → Misses table structures
- ✗ Figure captioning → Expensive specialized models
- ✗ Chunking → Loses visual context
- ✗ Embed text only → No visual information

Result: 40-60% information loss on visual documents



USING COMPLEX RETRIEVAL SYSTEMS THAT RELY ON OCR, DOCUMENT LAYOUT RECOGNITION, CHUNKING STRATEGIES, FIGURE CAPTIONING AND POWERFUL TEXT EMBEDDING MODELS

JUST EMBED THE IMAGE

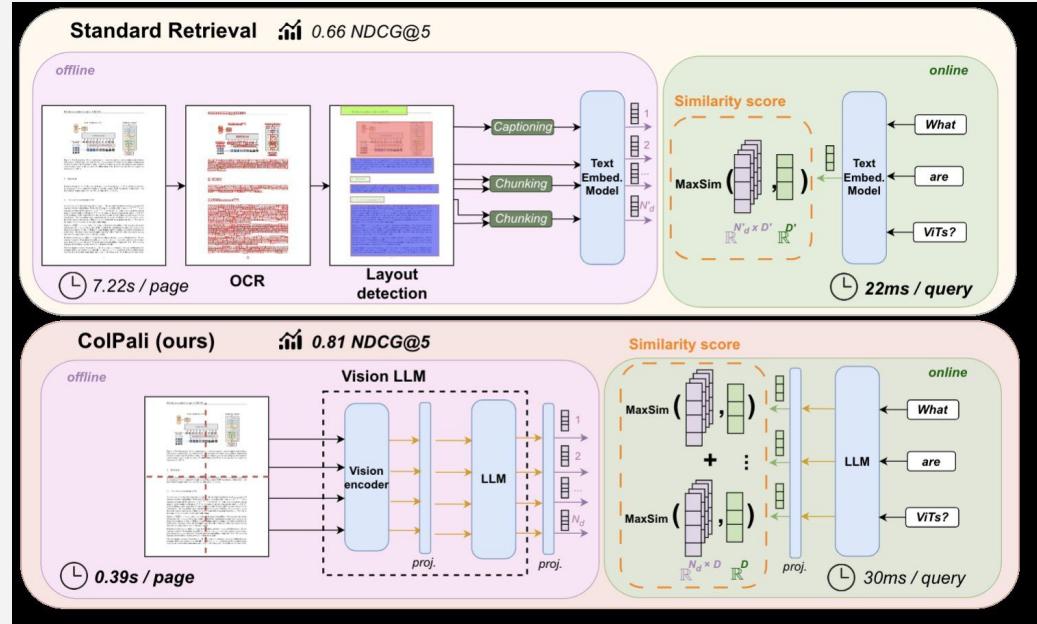
ColPali Revolution

Treat Pages as Images (No OCR!)

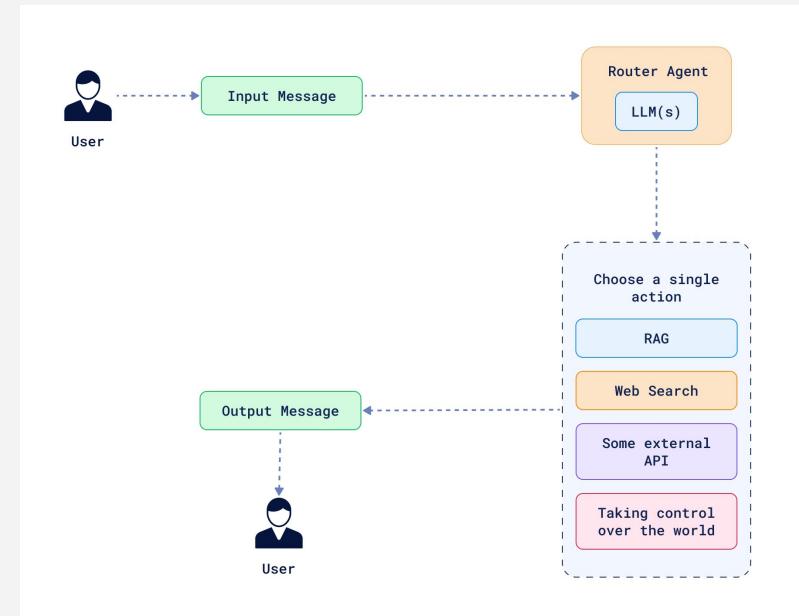
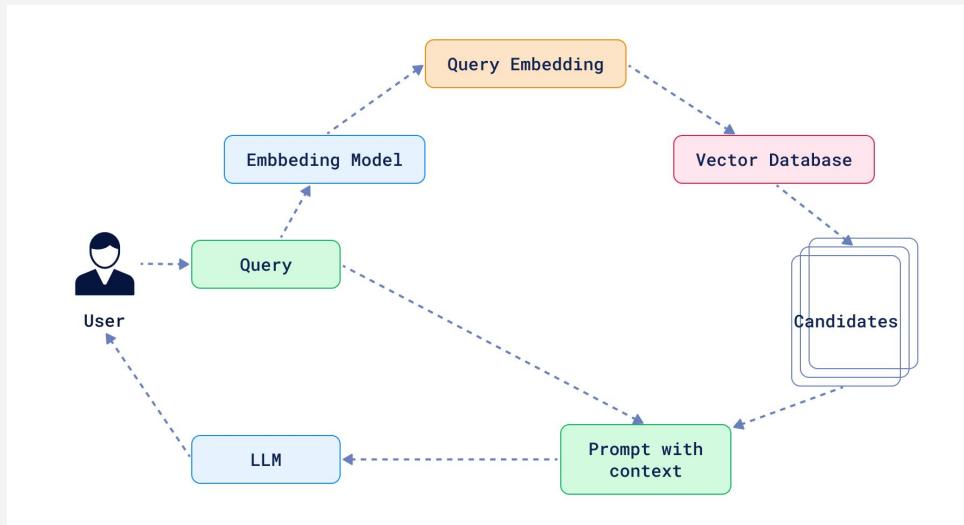


Architecture

- PaliGemma-3B Vision Language Model
- 32×32 patches = 1,024 patch embeddings per page
- ColBERT late interaction matching



What Makes RAG "Agentic"



Agentic Patterns



Query Routing

User: "What's the latest on AI regulation?"

Router: "latest" detected → Route to Web Search ✓

User: "What's our Q3 marketing budget?"

Router: "our" = internal → Route to Vector DB ✓

Query Decomposition

Complex: "Compare our Q3 to industry and predict Q4"

Agent breaks down:

- "Our Q3 metrics" → Internal DB
- "Industry Q3 benchmarks" → Web Search
- "Q4 factors" → Internal DB

Then synthesizes all results

Self-Correction

1. Retrieval
2. Grade docs
3. Score < threshold?
4. Rewrite query
5. Retry
6. Grade again

✓ **Use When:** Queries span sources, complexity varies

✗ **Don't Use:** Latency-critical (<500ms SLA), simple retrieval

Hybrid Search + Reranking Power



Vector-only: Misses exact matches, acronyms

BM25-only: No semantic understanding

Microsoft Azure AI Search Benchmarks:

- Vector-only: 43.8 NDCG@3
- Hybrid: 48.4 NDCG@3 (+10.5%)
- Hybrid + Reranker: 60.1 NDCG@3 (+37.2%)

Anthropic Contextual Retrieval Study:

- Baseline failure rate: 5.7%
- Hybrid + Reranker: 1.9%
- 67% reduction in retrieval failures

2 Stage Architecture Rerank:

Stage 1:

Similarity search → Top 100

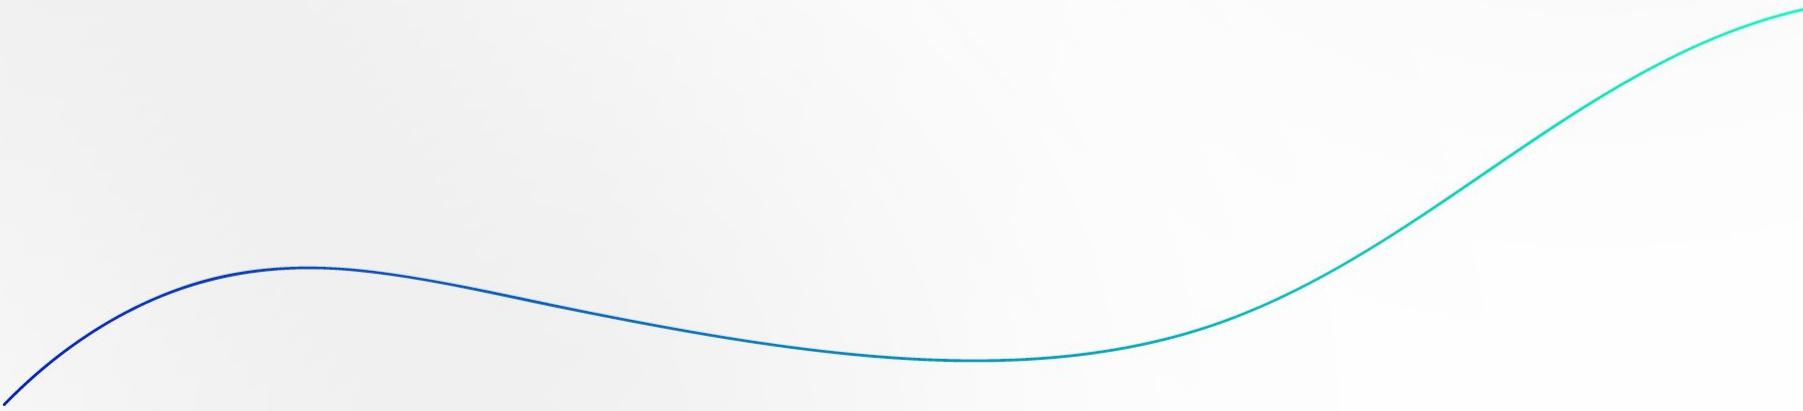
Stage 2:

Re-ranking pipeline → Top 5-10

✓ **Use When:** Queries span sources, complexity varies

✗ **Don't Use:** Latency-critical (<500ms SLA), simple retrieval

Production readiness



Production Best Practices



Do's

- 1. Use Hybrid Search (vector + BM25) as baseline
- 2. Implement metadata filtering (security-critical)
`filter = {"department": "finance", "access_level": user.role}`
- 3. Monitor retrieval quality (recall@k, NDCG)
- Tools: TruLens, LangSmith, DeepEval
- 4. Keep embeddings fresh (re-index on doc changes)
- 5. Evaluate systematically (not "vibes-based")



Don'ts

- 1. Vector-only search (use hybrid)
- 2. Ignore access controls (data leakage = lawsuit)
- 3. Overload context window (keep <50%)
- 4. Skip evaluation frameworks (doesn't scale)
- 5. Neglect data freshness (stale = wrong answers)





Thank you!

